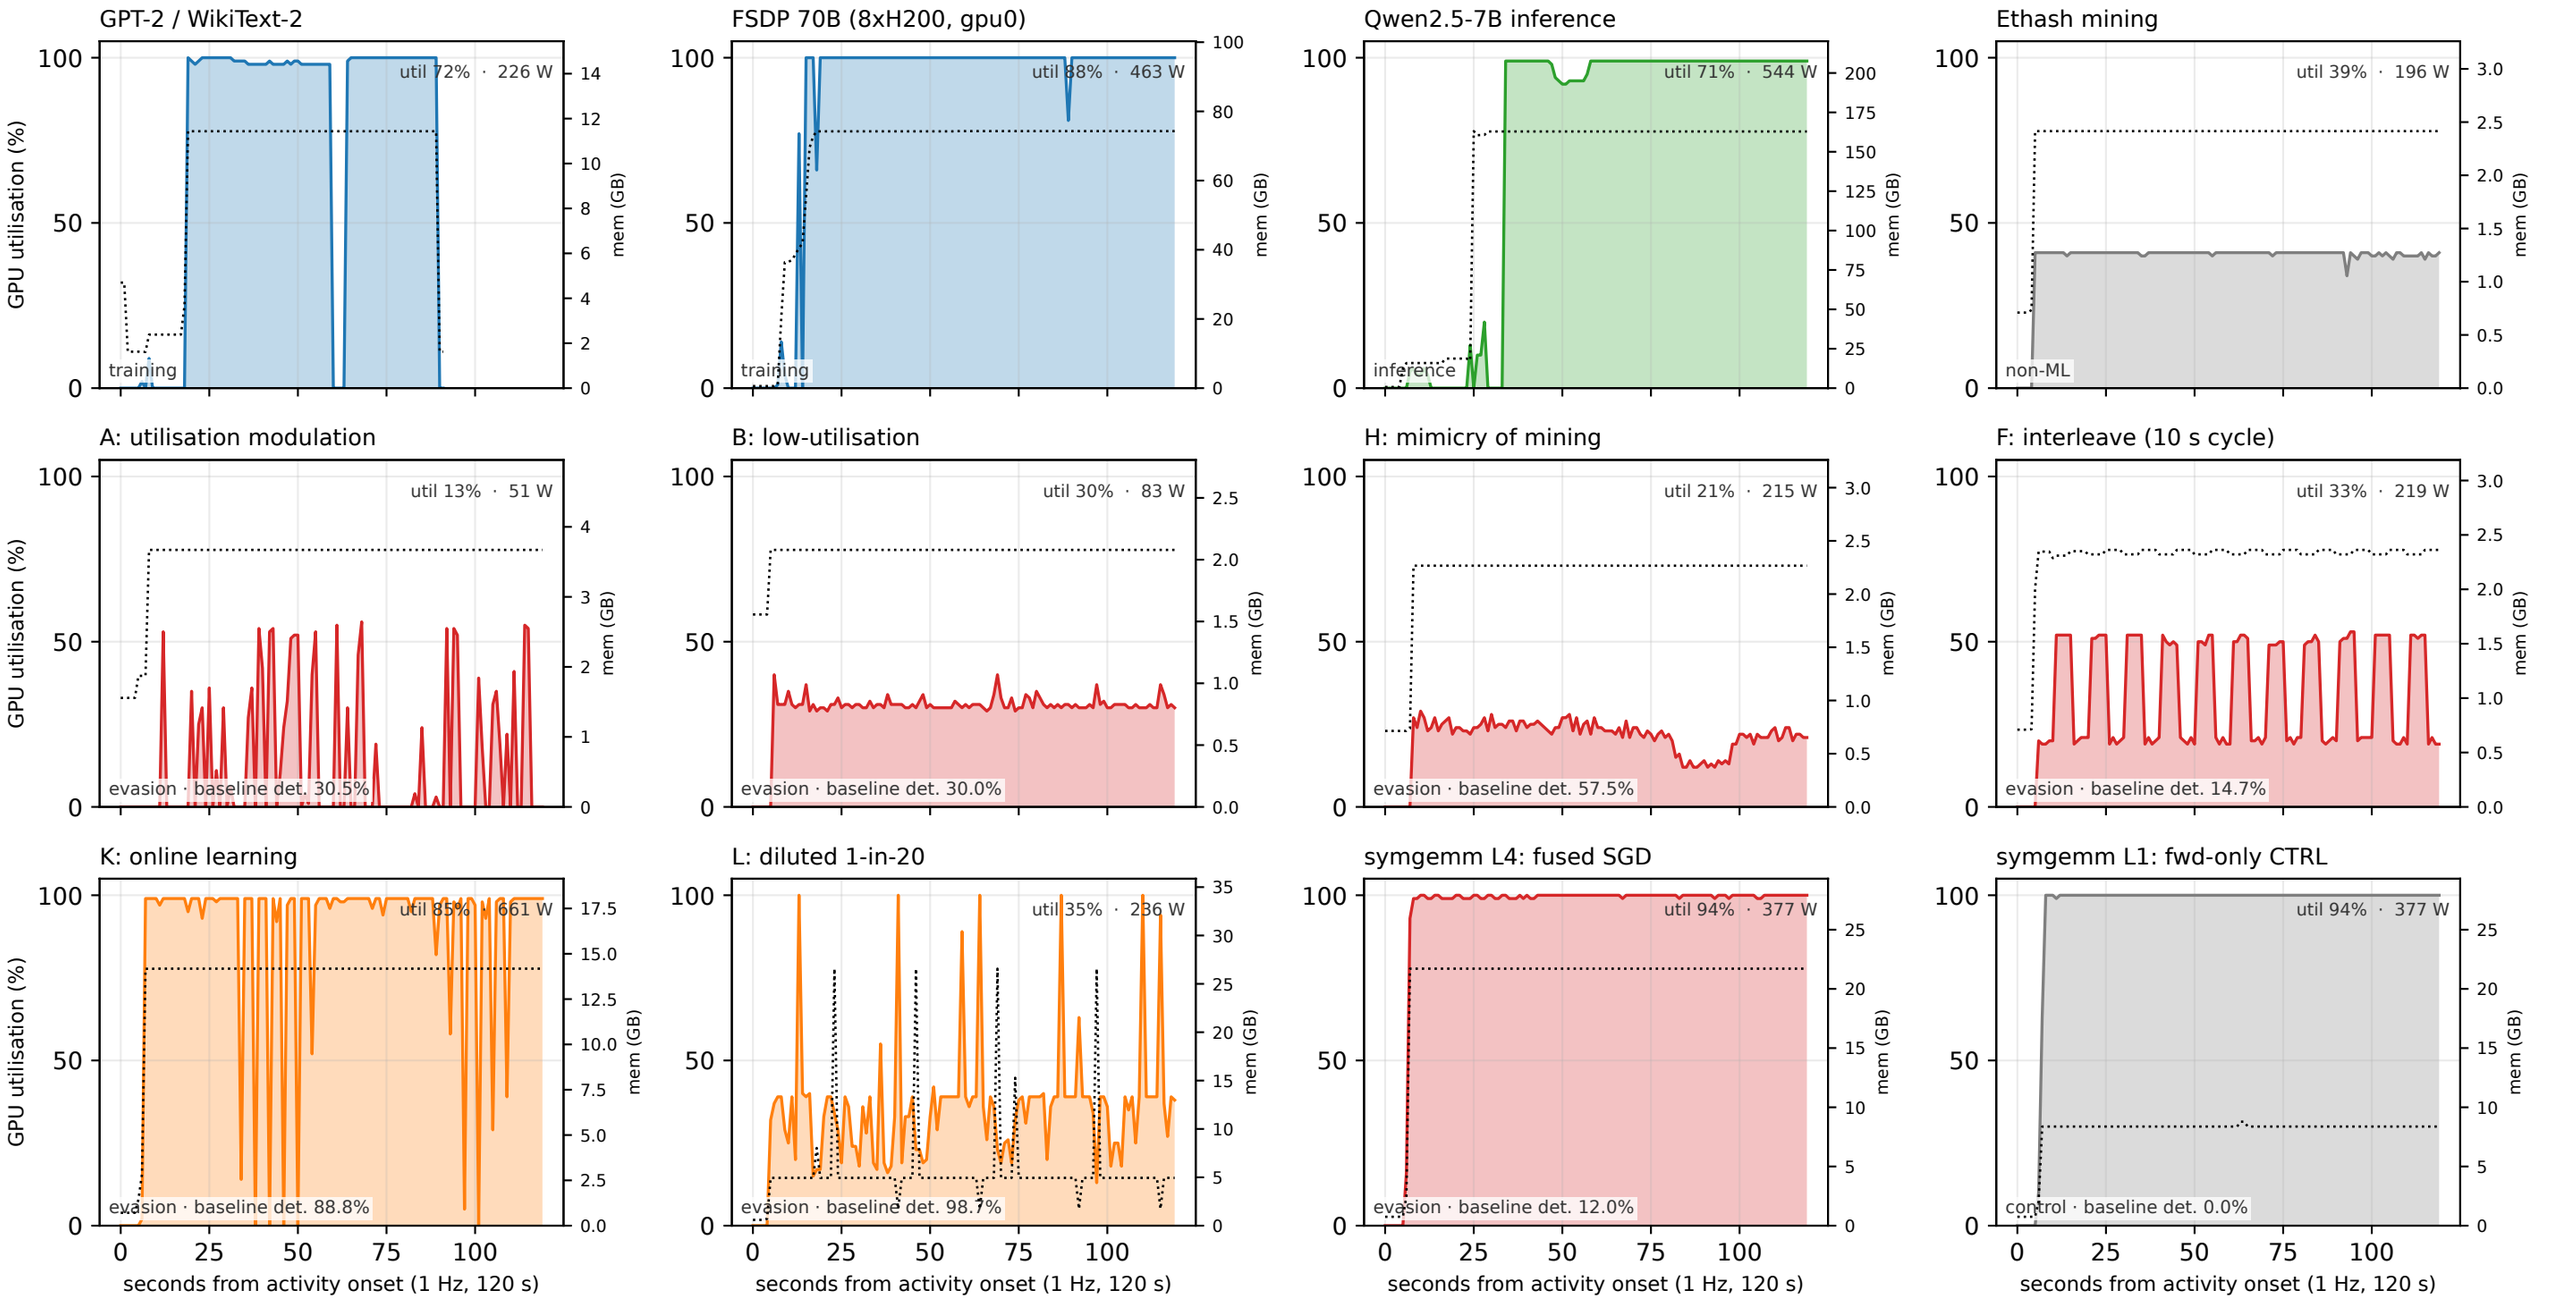
Raw NVML telemetry across the corpus — 120 s of one representative run per workload, aligned on activity onset (shaded: utilisation; dotted: memory)
GPT-2 / WikiText-2
util 72% · 226 W
training
FSDP 70B (8xH200, gpu0)
util 88% · 463 W
training
Qwen2.5-7B inference
util 71% · 544 W
inference
Ethash mining
util 39% · 196 W
non-ML
A: utilisation modulation
util 13% · 51 W
evasion · baseline det. 30.5%
B: low-utilisation
util 30% · 83 W
evasion · baseline det. 30.0%
H: mimicry of mining
util 21% · 215 W
evasion · baseline det. 57.5%
F: interleave (10 s cycle)
util 33% · 219 W
evasion · baseline det. 14.7%
K: online learning
util 85% · 661 W
evasion · baseline det. 88.8%
L: diluted 1-in-20
util 35% · 236 W
evasion · baseline det. 98.7%
symgemm L4: fused SGD
util 94% · 377 W
evasion · baseline det. 12.0%
symgemm L1: fwd-only CTRL
util 94% · 377 W
control · baseline det. 0.0%
GPU utilisation (%)
mem (GB)
seconds from activity onset (1 Hz, 120 s)