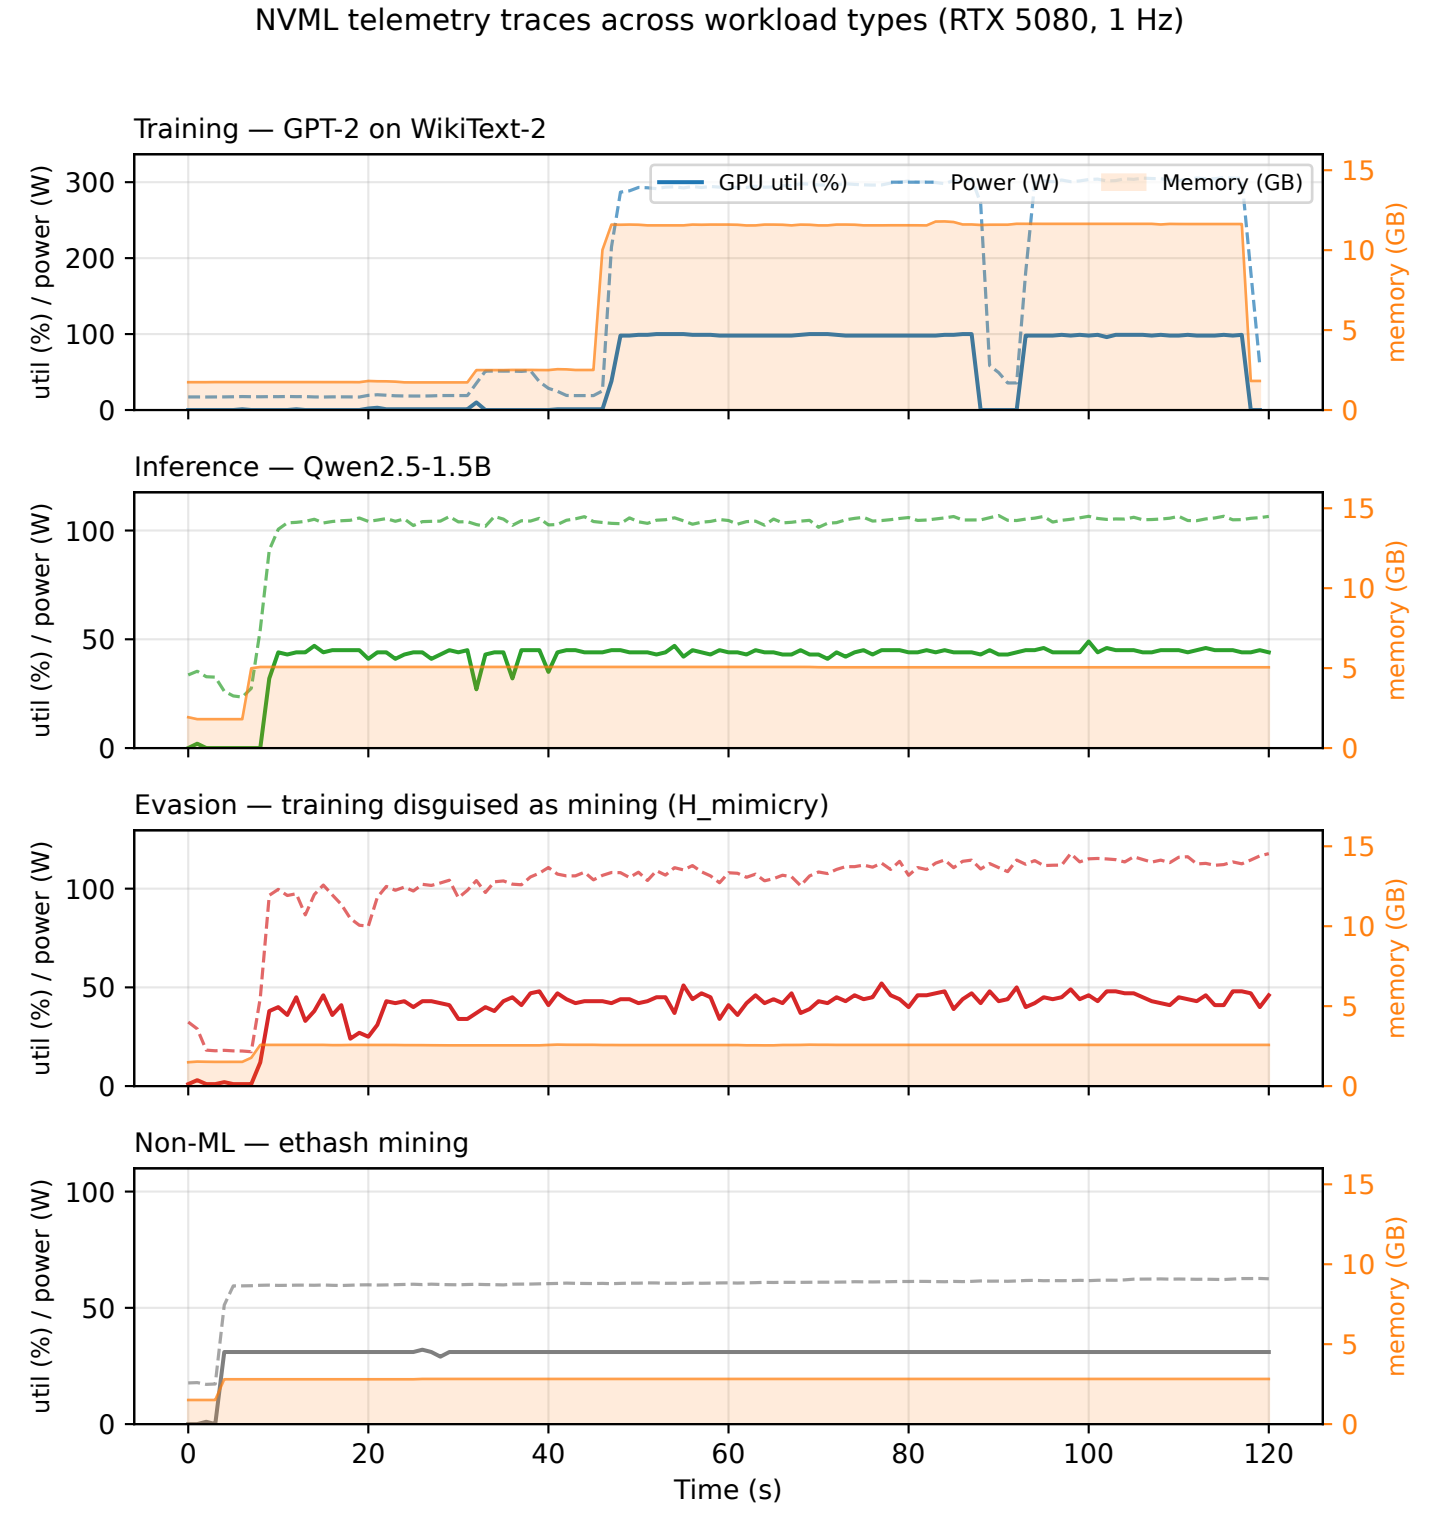
NVML telemetry traces across workload types (RTX 5080, 1 Hz)
Training — GPT-2 on WikiText-2
GPU util (%)
Power (W)
Memory (GB)
util (%) / power (W)
memory (GB)
Inference — Qwen2.5-1.5B
Evasion — training disguised as mining (H_mimicry)
Non-ML — ethash mining
Time (s)